



FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Weakly and Self-Supervised Learning

**A. Maier, V. Christlein, K. Breininger, S. Vesal, F. Meister, C. Liu, S. Gündel, S. Jaganathan, N. Maul,
M. Vornehm, L. Reeb, F. Thamm, M. Hoffmann, C. Bergler, F. Denzinger, W. Fu, B. Geissler, Z. Yang**
Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg
July 7, 2020



Outline

Learning with Limited Annotations

Definition

Image-based SSL for Representation Learning

Generative

Spatial Context

Context Similarity

Contrastive SSL

Supervised Contrastive Learning

Bootstrap SSL – A paradigm change



FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Learning with Limited Annotations



Supervised Learning

So far, we have seen impressive results, achieved with ...

- ... large amounts of training data and
- ... consistent, high-quality annotations



Mask R-CNN [7], image source [2]

The Cost of Annotation

Image-level class labels: ~27 sec [11]



Instance spotting: + 14 sec [11]



Instance Segmentation: + 80 sec [11]

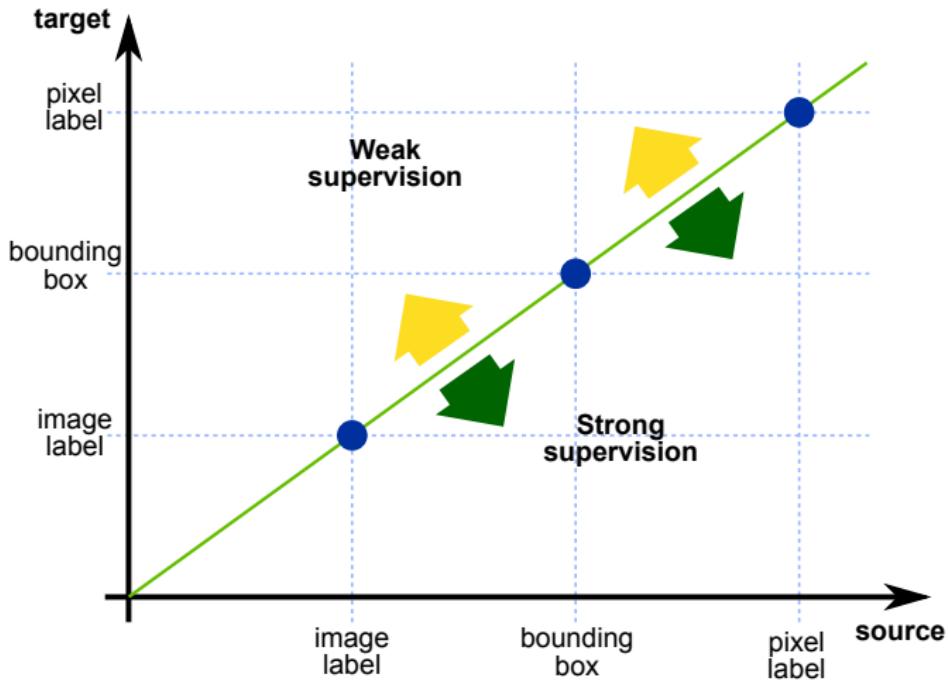


Dense pixel-level annotations: 1.5h [4]



Source: [4], [11]

Strongly vs. Weakly Supervised Learning



Source: Reproduced from CVPR18 Tutorial: Weakly Supervised Learning for Computer Vision

Key Ingredients for Weakly Supervised Learning

Priors: Explicit and Implicit

- Shape + size
- Contrast
- Motion
- Class distribution
- Similarity across images

Hints

- Image labels
- Bounding boxes
- Image caption
- Sparse temporal labels
- Scribbles
- Clicks inside objects

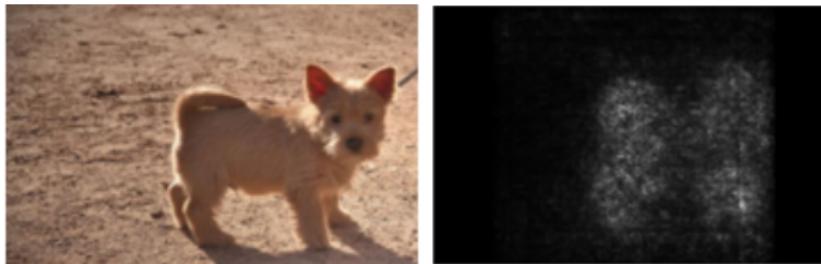


Source: Bearman et al. [3]

From Labels to Localization

Approach 1: Use a **pretrained** classification network [13]

- How does a change in the input affect the classification?
→ Lecture on Visualization
- Qualitative segmentation map
- **Problem 1:** Classifier was never trained for localized decisions
- **Problem 2:** Good classifiers don't automatically yield good maps



Source: Simonyan et al. [13]

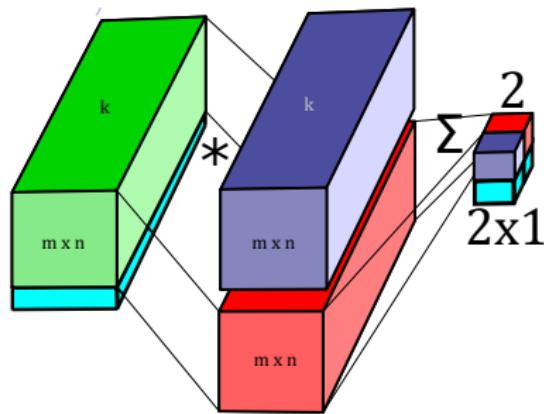
From Labels to Localization

Approach 2: Use a classification network, but smarter [14]

- Core idea: Use **global average pooling**
→ Fully convolutional networks revisited

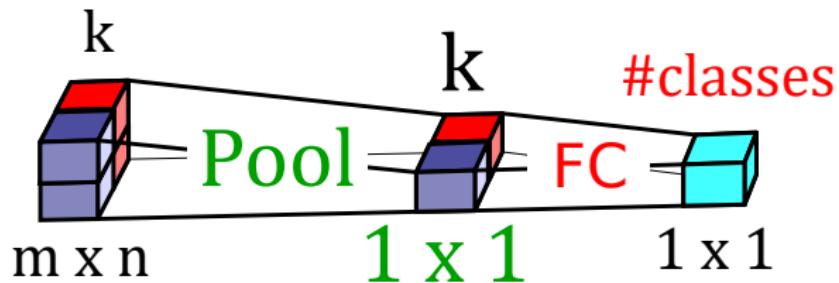
Fully Convolutional Networks: Revisited

- Fully connected layers fix the input size
→ replace by $m \times n$ convolution
- Alternatively, we can also **pool** to the correct size first



Fully Convolutional Networks: Revisited

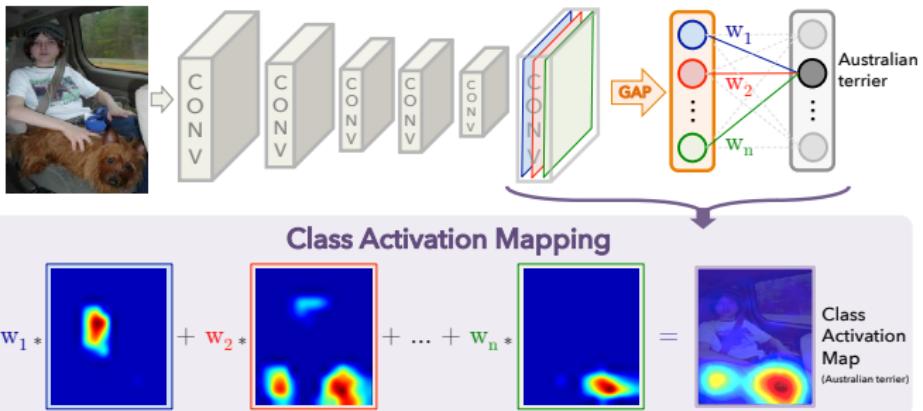
- Fully connected layers fix the input size
→ replace by $m \times n$ convolution
- Alternatively, we can also **pool** to the correct size first



From Labels to Localization

Approach 2: Use a classification network, but smarter [14]

- Core idea: Use **global average pooling**
- Then look at penultimate layer
- **Class Activation Maps (CAMs)**
- Generalization: Grad-CAM [12]



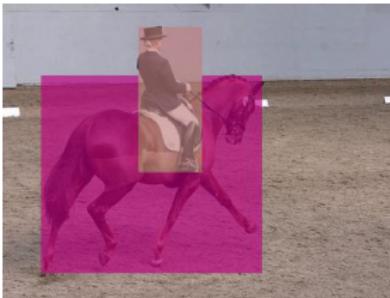
Source: Zhou et al. [14]

From Bounding Boxes to Segmentation

Expensively annotated
Fully supervised



Cheaply annotated



Cheaply annotated
Weakly supervised

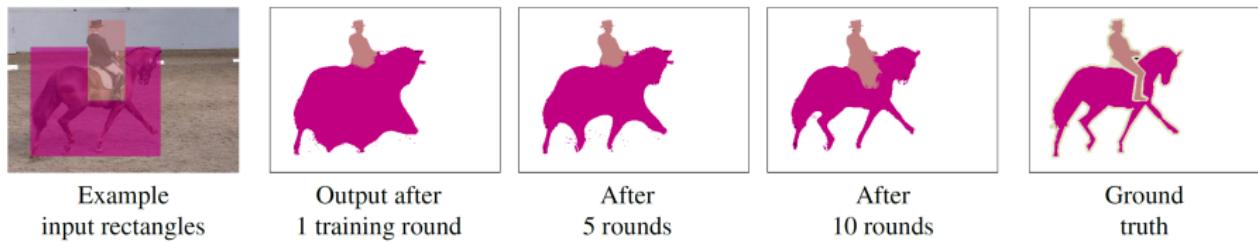


- Manual segmentation is tedious
- Bounding boxes are less tedious
- Can we learn segmentation from boxes?

Source: Khoreva et al. [6]

From Bounding Boxes to Segmentation

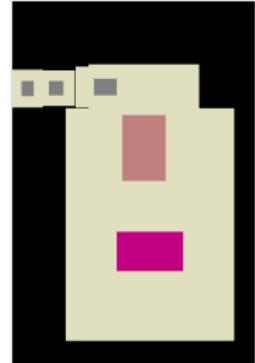
- Observation: Convolutional NNs are somewhat robust to (label-)noise
- Let's use that: Use bounding boxes as target and recursively estimate better targets [6]



- Problem: Training quickly degrades
- **Postprocess** intermediate predictions

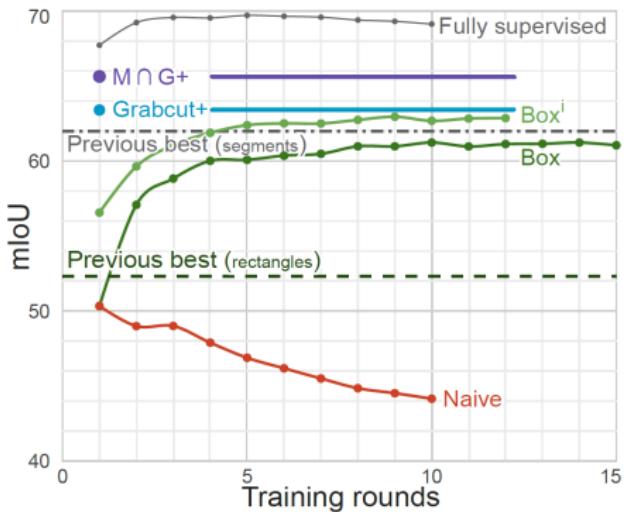
From Bounding Boxes to Segmentation

- **SUPPRESS** detections
 - ...of wrong class
 - ...outside the box
 - ... $<\%$ box area
 - ...outside of conditional random field boundaries
- Additional improvement: smaller boxes
 - Objects are “on average” roundish
 - Corners and edges contain “on average” the least true positives
 - Define “ignore” regions with unknown labels



Source: Khoreva et al. [6]

Improved Recursive Training



- Shrinking boxes beats state of the art
- Combine Grabcut and MCG for initial label
 - No need for recursion

Source: [6]

NEXT TIME
ON DEEP LEARNING



FAU

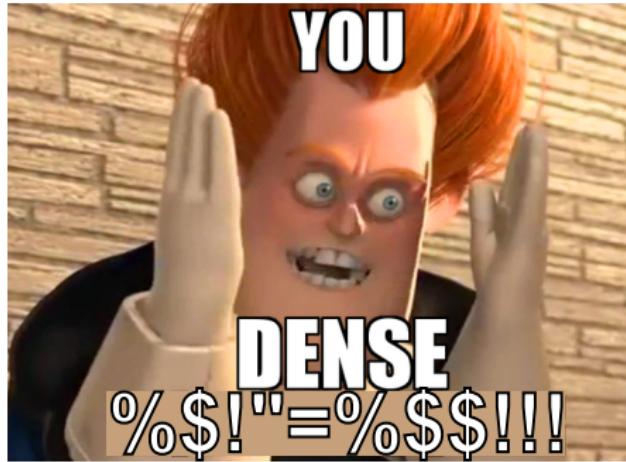
FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Weakly and Self-Supervised Learning - Part 2

**A. Maier, V. Christlein, K. Breininger, S. Vesal, F. Meister, C. Liu, S. Gündel, S. Jaganathan, N. Maul,
M. Vornehm, L. Reeb, F. Thamm, M. Hoffmann, C. Bergler, F. Denzinger, W. Fu, B. Geissler, Z. Yang**
Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg
July 7, 2020



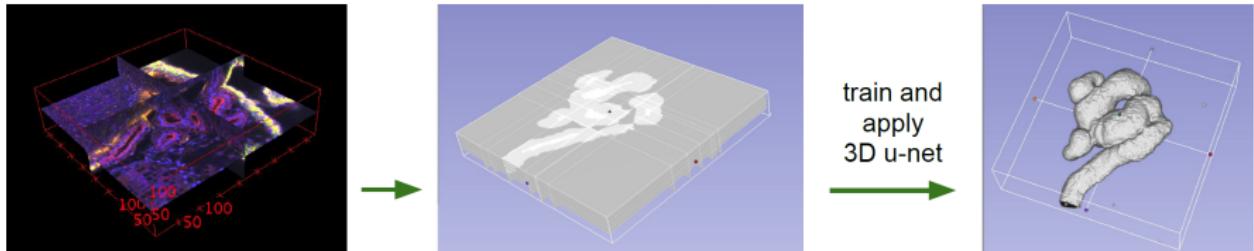
From Sparse Annotations to Dense Segmentations



... not quite

Source: Adapted from <https://knowyourmeme.com/memes/>.

From Sparse Annotations to Dense Segmentations



- 3D segmentation is extremely tedious
- Obtain only a few labelled 2D slides
- Compute automatic segmentation in 3D
- Allows for interactive correction

Source: Çiçek et al. [1]

From Sparse Annotations to Dense Segmentations

Training with sparse labels

- Problem: “One hot” labels $y_{n,i}$ with element i being 1
 → either **true** or **false**

$$L(\mathbf{y}, \hat{\mathbf{y}}) = \sum_n -\log \hat{y}_{n,i} \Big|_{y_{n,i}=1}$$

- Obtain sparse loss by weighted cross entropy [1]

$$L'(\mathbf{y}, \hat{\mathbf{y}}) = L(\mathbf{y}, \hat{\mathbf{y}}) \cdot w(\mathbf{y})$$

where $w(y_{n,i}) = \begin{cases} 0 & \text{if } y_n \text{ is not labelled} \\ w_{n,i} > 0 & \text{otherwise} \end{cases}$

- Can be easily extended to interactive segmentation by updating \mathbf{y} and $w(\mathbf{y})$

Take Away: Weakly Supervised Learning

- Fine grained labels are expensive - can we get away with something cheaper?
- Core definition: **Label** has less detail than **target**
- Methods depend on **prior knowledge** and **weak labels** ("hints")
- Typically inferior to **fully supervised** training
 - but highly relevant in practice
- Don't forget transfer learning (!)
- Related:
 - **Semi-supervised** Learning
 - **Self-supervised** Learning

**NEXT TIME
ON DEEP LEARNING**



FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Weakly and Self-Supervised Learning - Part 3

**A. Maier, V. Christlein, K. Breininger, S. Vesal, F. Meister, C. Liu, S. Gündel, S. Jaganathan, N. Maul,
M. Vornehm, L. Reeb, F. Thamm, M. Hoffmann, C. Bergler, F. Denzinger, W. Fu, B. Geissler, Z. Yang**
Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg
July 7, 2020



Outline

Learning with Limited Annotations

Definition

Image-based SSL for Representation Learning

Generative

Spatial Context

Context Similarity

Contrastive SSL

Supervised Contrastive Learning

Bootstrap SSL – A paradigm change



FAU

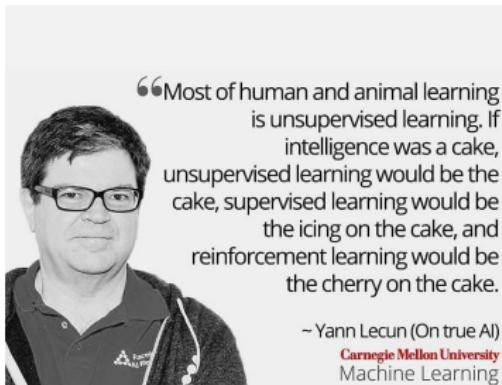
FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Definition



Motivation

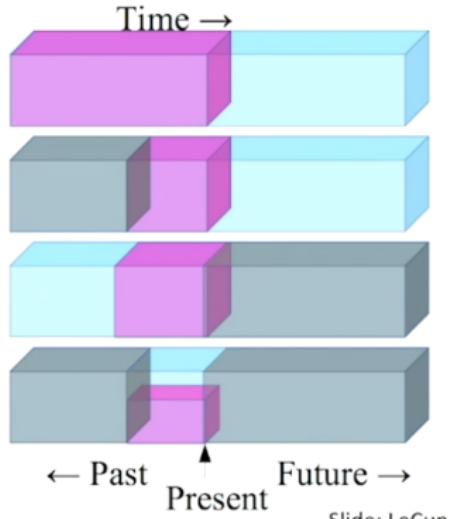
- Jitendra Malik: “Supervision is the opium of the AI researcher”
- Alyosha Efros: “The AI revolution will not be supervised”
- Yann LeCun:



Source: <https://www.facebook.com/722677142/posts/10156036317282143/>

Idea

- ▶ Predict any part of the input from any other part.
- ▶ Predict the **future** from the **past**.
- ▶ Predict the **future** from the **recent past**.
- ▶ Predict the **past** from the **present**.
- ▶ Predict the **top** from the **bottom**.
- ▶ Predict the **occluded** from the **visible**
- ▶ **Pretend there is a part of the input you don't know and predict that.**



Slide: LeCun

Source: <https://www.youtube.com/watch?v=7I0Qt7GALVk>

Self-supervised Learning: Definition



Yann LeCun

April 30, 2019 · ②

I now call it "self-supervised learning", because
 "unsupervised" is both a loaded and confusing term.

- Subcategory of unsupervised learning
- Use pretext/surrogate/pseudo tasks in a supervised fashion
 - Automatically generated labels
 - Measurement of correctness
- Downstream task: retrieval, supervised or semi-supervised classification, etc.

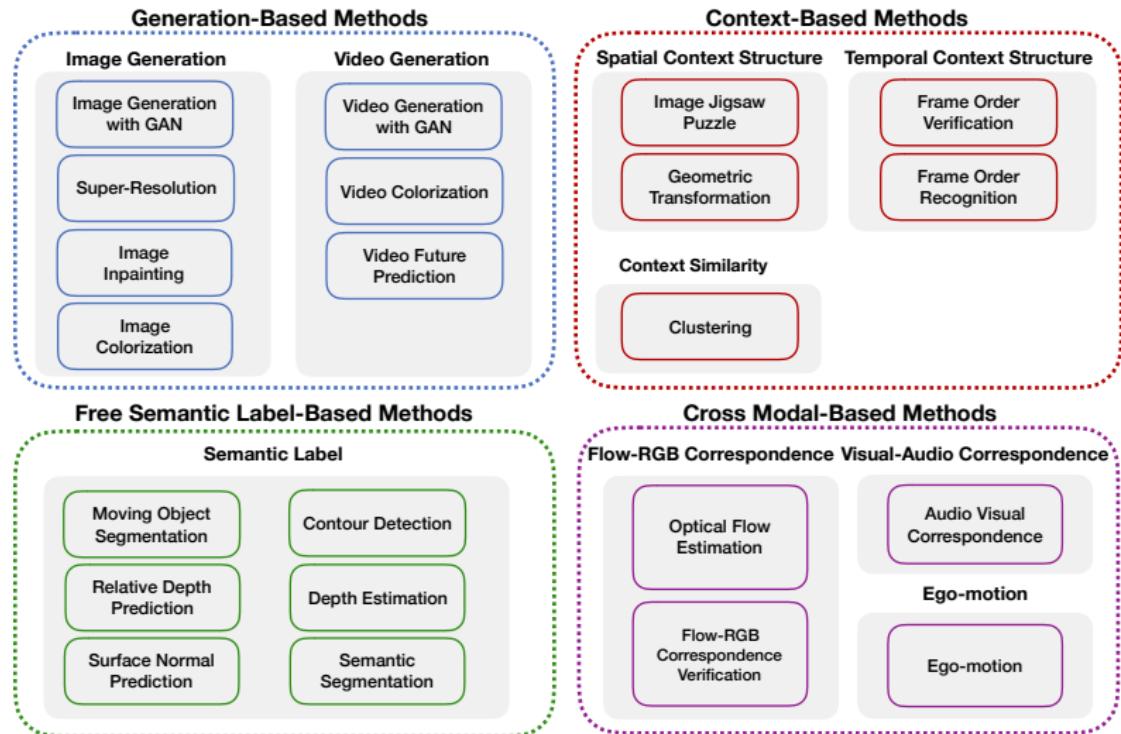
Note: Generative models (e.g. GANs) are also SSL methods

Source: <https://www.facebook.com/722677142/posts/10155934004262143/>



Source: [15]

Pretext Tasks Overview



Source: [15]



FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

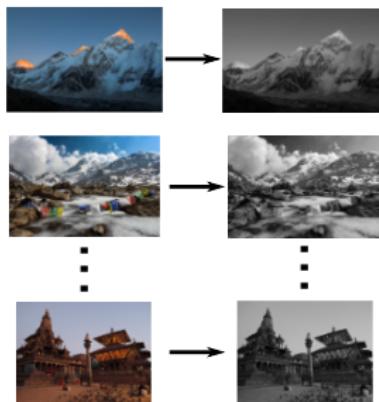
Image-based SSL for Representation Learning



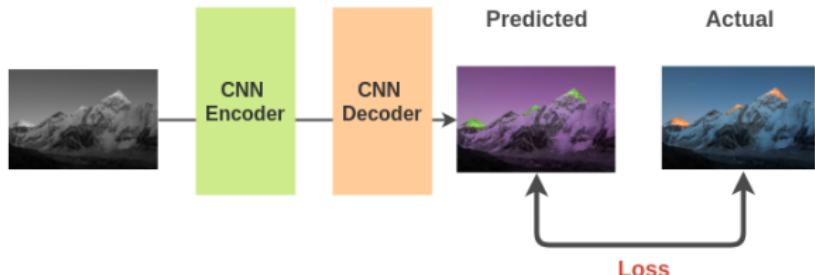
Generative

Image Colorization

Data generation:



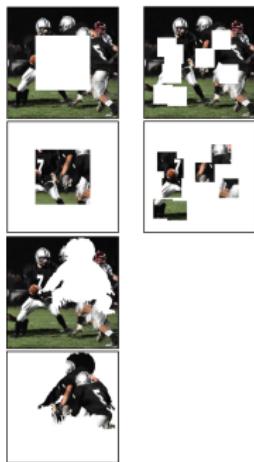
Pretext task: ℓ_2 loss between gray and color version



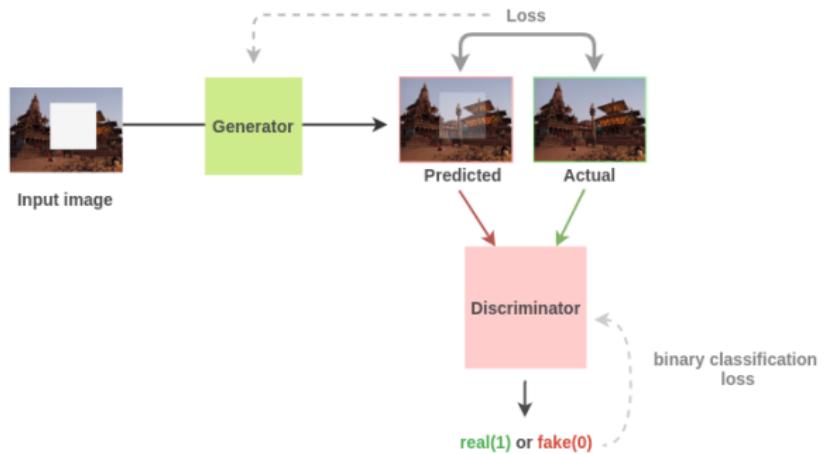
Source: <https://amitness.com/2020/02/illustrated-self-supervised-learning/>

Image Inpainting

Data generation:



Pretext task:

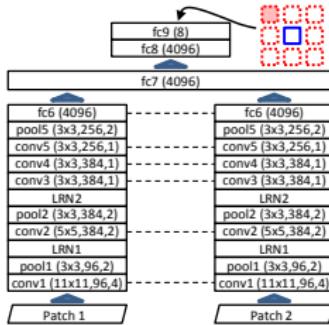
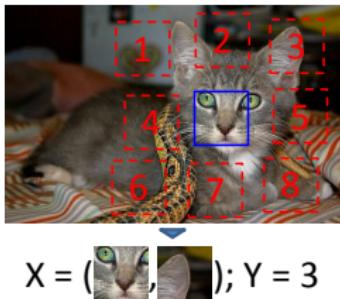


Source: [16]

Source: <https://amitness.com/2020/02/illustrated-self-supervised-learning/>

Spatial Context

Solve Jigsaw Puzzle [17]

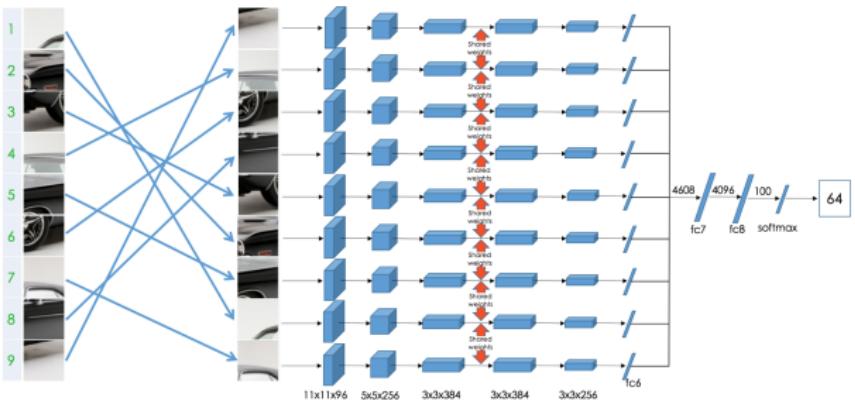
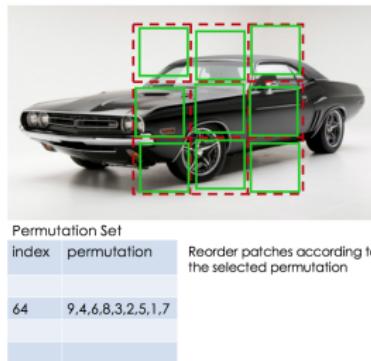


Attention: Trivial solution possible

- boundary patterns, continuing textures → use large enough gaps
- chromatic aberration
 - Pre-process images by shifting green and magenta toward gray
 - randomly drop 2 color channels

Source: [17]

Solve Jigsaw Puzzle++ [18]



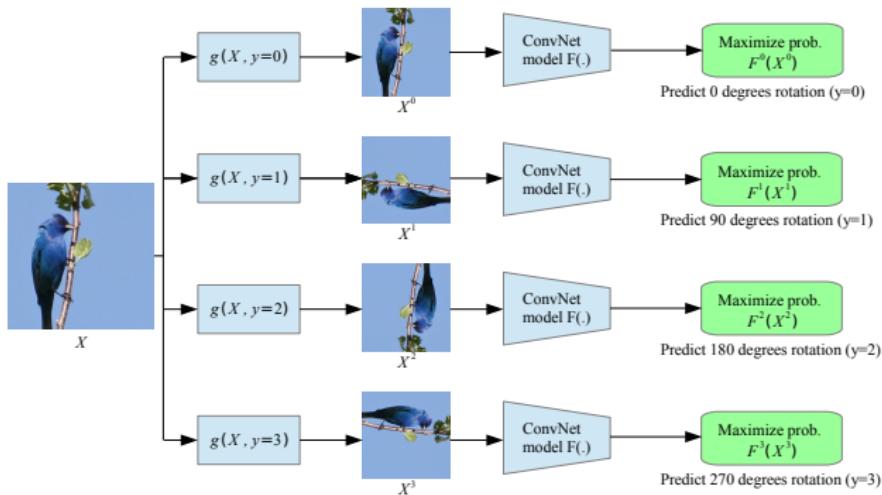
9 tiles $\rightarrow 9! = 362\,880$ possible permutations

Source: [18]

Solve Jigsaw Puzzle++ [18] (cont.)

Number of permutations	Average hamming distance	Minimum hamming distance	Jigsaw task accuracy	Detection performance
1000	8.00	2	71	53.2
1000	6.35	2	62	51.3
1000	3.99	2	54	50.2
100	8.08	2	88	52.6
95	8.08	3	90	52.4
85	8.07	4	91	52.7
71	8.07	5	92	52.8
35	8.13	6	94	52.6
10	8.57	7	97	49.2
7	8.95	8	98	49.6
6	9	9	99	49.7

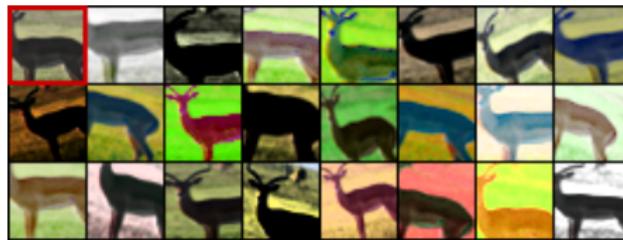
Rotation [19]



Source: [19]

Context Similarity

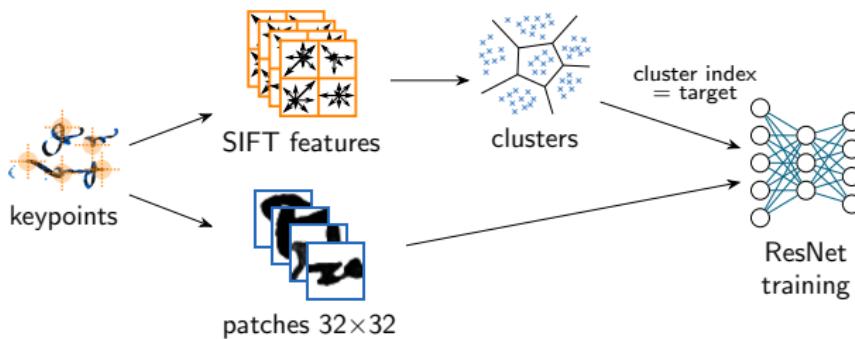
Distortions [21] (Exemplar-CNN)



- For each input patch, create N (e.g. $N = 100$) distorted images
 - All these distorted images form one class
- Discriminate between a set of surrogate classes (e.g. 8000 pseudo-classes)

Source: [21]

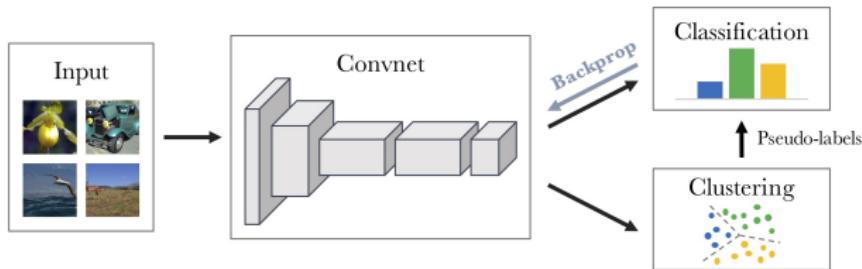
Clustering [22]



- From keypoints in an image extract patches and compute descriptors
- Cluster features of patches using k -means into N clusters ($N = 5000$)
- Use cluster indices as targets for input patches
- Remove features (+patches) in between of two clusters

Source: [22]

Clustering [20] (DeepCluster)

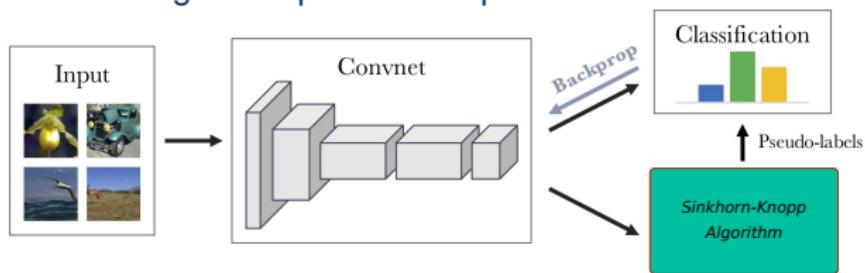


- Alternate between
 - k -means clustering (each epoch) of PCA-whitened ($D = 256$) & ℓ_2 -normalized activation features
 - CNN training
- Avoid trivial solutions
 - Re-assign empty clusters
 - Weight contribution of an input by inverse of the size of its assigned cluster

Source: [20]

Clustering [24]

Self-labelling with Optimal Transport



Optimal transport

Supply		Need	
25 laptops	Warehouse A	Shop 1	25 laptops
25 laptops	Warehouse B	Shop 2	25 laptops

Distance(cost) matrix → Optimal Allocation

	Shop 1	Shop 2	
Warehouse A	2km	3km	Shop 1
Warehouse B	2km	1km	Shop 2
	25	0	
	0	25	

Source: <https://amitness.com/2020/04/illustrated-self-labelling/>

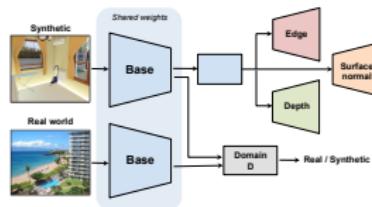
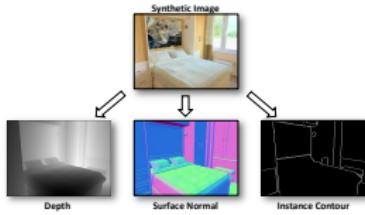
Clustering [24] (cont.)

Self-labelling with Optimal Transport

Comparison to DeepCluster

- no separate clustering loss → can lead to degenerate solutions
- clustering approach that minimizes the same cross-entropy loss that the network also optimize.

Multi-task SSL using Synthetic Imagery [23]



- Given: input synthetic RGB image
- Network simultaneously predicts: surface normal, depth, instance contour
- Additionally: minimize feature space domain differences between real and synthetic data

Source: [23]

**NEXT TIME
ON DEEP LEARNING**



FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Weakly and Self-Supervised Learning - Part 4

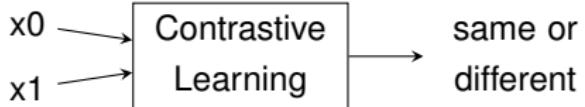
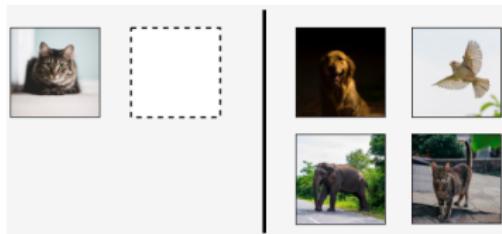
**A. Maier, V. Christlein, K. Breininger, S. Vesal, F. Meister, C. Liu, S. Gündel, S. Jaganathan, N. Maul,
M. Vornehm, L. Reeb, F. Thamm, M. Hoffmann, C. Bergler, F. Denzinger, W. Fu, B. Geissler, Z. Yang**
Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg
July 7, 2020



Contrastive SSL

Contrastive Learning

Match the correct animal



Advantages over generative/context models:

- Pixel-level losses could overly focus on pixel-based details, rather than more abstract latent factors
- Pixel-based objectives often assume pixel independence → reduce ability to model correlations or complex structure

Source: <https://amitness.com/2020/03/illustrated-simclr/>

Contrastive Loss

Given: $\mathcal{X} = \{\mathbf{x}, \underbrace{\mathbf{x}^+}_{\text{positive sample}}, \underbrace{\mathbf{x}_1^-, \dots, \mathbf{x}_{N-1}^-}_{\text{negative samples}}\}$; similarity function $s(\cdot)$ (e.g. cosine similarity)

Goal: $s(f(\mathbf{x}), f(\mathbf{x}^+)) >> s(f(\mathbf{x}), f(\mathbf{x}^-))$

Contrastive/InfoNCE Loss (aka ‘n-pair loss’/‘consistency loss’/‘ranking-based NCE’):

Cross-entropy loss for (N)-way softmax classifier

$$\begin{aligned}\mathcal{L}_N &= -\mathbb{E}_{\mathcal{X}} \left[\log \frac{\exp(s(f(\mathbf{x}), f(\mathbf{x}^+)))}{\exp(s(f(\mathbf{x}), f(\mathbf{x}^+))) + \sum_{j=1}^{N-1} \exp(s(f(\mathbf{x}), f(\mathbf{x}_j^-)))} \right] \\ &= -\mathbb{E}_{\mathcal{X}} \left[\log \frac{\exp(s(f(\mathbf{x}), f(\mathbf{x}^+)))}{\sum_{j=1}^N \exp(s(f(\mathbf{x}), f(\mathbf{x}_j)))} \right]\end{aligned}$$

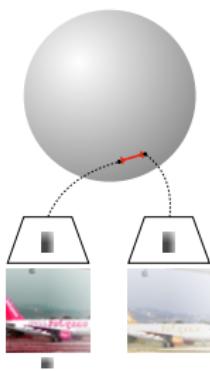
Contrastive Loss (cont.)

Minimizing Contrastive Loss maximizes a lower bound on the mutual information between $f(\mathbf{x})$ and $f(\mathbf{x}^+)$ [25], [27].

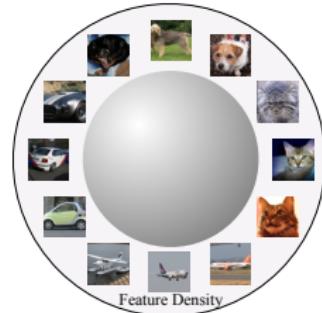
Common Variation: temperature hyperparameter τ

$$\mathcal{L}_N = -\mathbb{E}_{\mathcal{X}} \left[\log \frac{\exp(s(f(\mathbf{x}), f(\mathbf{x}^+))/\tau)}{\sum_{j=1}^{N+1} \exp(s(f(\mathbf{x}), f(\mathbf{x}_j))/\tau)} \right]$$

Effectivity of Contrastive Loss



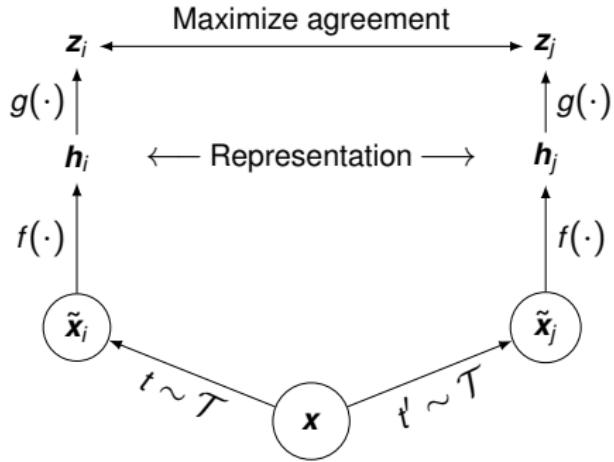
Alignment: Similar samples have similar features



Uniformity: Preserve maximal information.

Source: [35]

Examples: SimCLR [31]



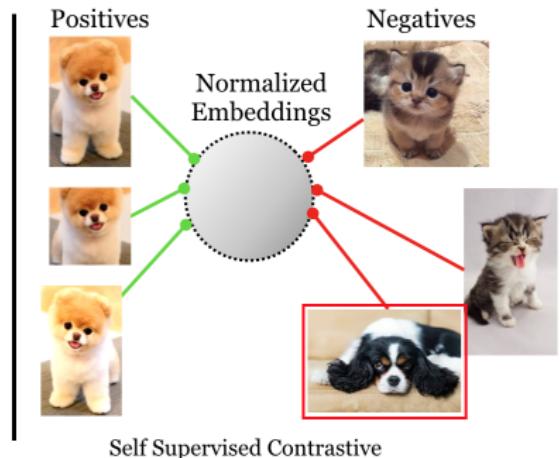
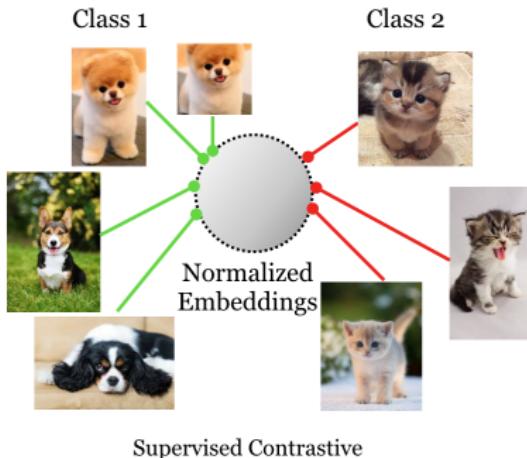
1. Mini-batch of n samples. Each sample is applied with two different data augmentation operations $\rightarrow 2n$ augmented samples: $\tilde{\mathbf{x}}_i = f(\mathbf{x}), \quad \tilde{\mathbf{x}}_j = f'(\mathbf{x}), \quad t, t' \sim \mathcal{T}$
2. One positive pair, $2(n - 1)$ negatives. Representation through base encoder f :

$$\mathbf{h}_i = f(\tilde{\mathbf{x}}_i), \quad \mathbf{h}_j = f(\tilde{\mathbf{x}}_j)$$
3. Contrastive loss on top of representation head g :

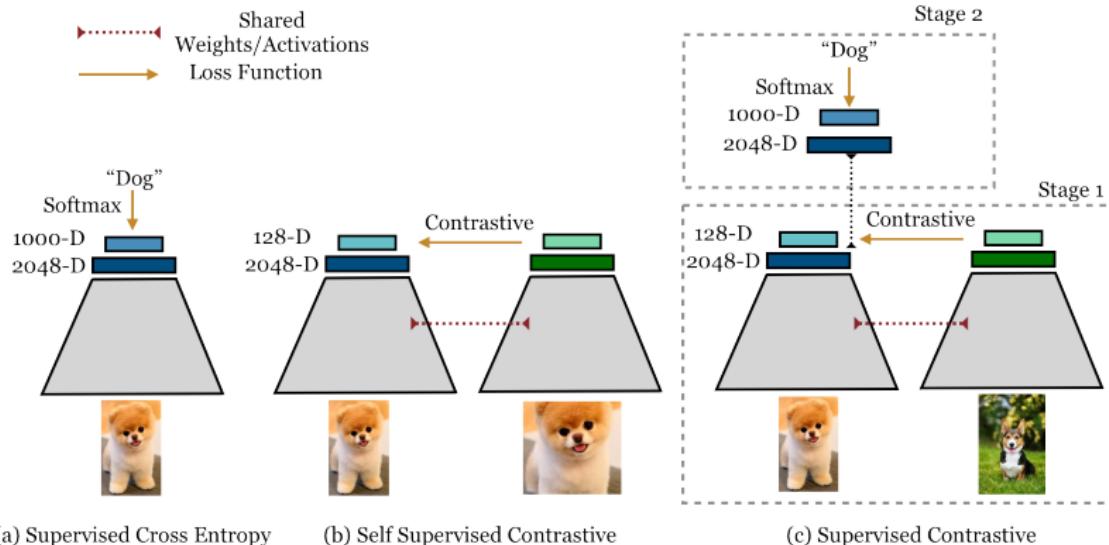
$$\mathcal{L}_{i,j} = -\log \frac{\exp(s(\mathbf{z}_i, \mathbf{z}_j) / \tau)}{\sum_{k=1}^{2n} \mathbf{1}_{[k \neq i]} \exp(s(\mathbf{z}_i, \mathbf{z}_k) / \tau)}$$

Supervised Contrastive Learning

Supervised Contrastive Learning [33]



Supervised Contrastive Learning [33]



Supervised Contrastive Loss

- Self-supervised **has no** knowledge about class labels → only one positive example
- Supervised **has** knowledge about class labels → many positives per example
- Compute loss between any sample \mathbf{z}_j having the same class as anchor \mathbf{z}_i ($\mathbf{y}_i = \mathbf{y}_j$)

$$L_{\text{sup}} = \sum_{i=1}^{2N} - \dots \sum_{j=1}^{2N} \mathbb{1}_{i \neq j} \cdot \mathbb{1}_{\mathbf{y}_i = \mathbf{y}_j} \cdot \log \frac{\exp(\mathbf{z}_i^\top \mathbf{z}_j / \tau)}{\sum_{k=1}^{2N} \mathbb{1}_{i \neq k} \cdot \exp(\mathbf{z}_i^\top \mathbf{z}_k / \tau)}$$

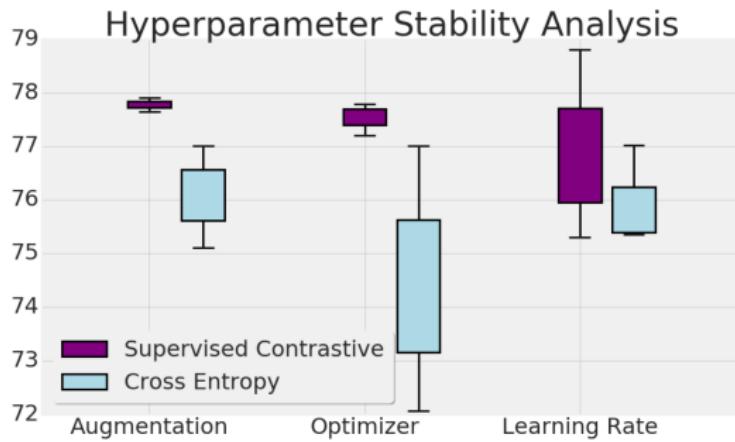
Supervised Contrastive Loss (cont.)

- Vectors \mathbf{z} need to be normalized, i.e. $\mathbf{z} = \mathbf{w}/\|\mathbf{w}\|$, where \mathbf{w} is the output of the projection network
- Gradient w.r.t. to \mathbf{w} is high for hard positives and negatives and small otherwise → “built-in” focal loss
- For one positive and one negative it turns out that

$$L_{\text{sup}} \propto \|\mathbf{z}_a - \mathbf{z}_p\|^2 - \|\mathbf{z}_a - \mathbf{z}_n\|^2 + 2\tau$$

→ Common contrastive loss in siamese networks

Hyperparameter stability



- Increased stability w.r.t. to non-optimal hyperparameters

Source: [33]

What else?

- Training about 50% slower than CE
- Improves over training with SOTA data augmentation (CutMix)
- Enables unsupervised clustering in latent space → correction of label-noise, new possibilities for semi-supervised, ...

Bootstrap SSL – A paradigm change

Bootstrap Your Own Latent (BYOL) [34] Overview

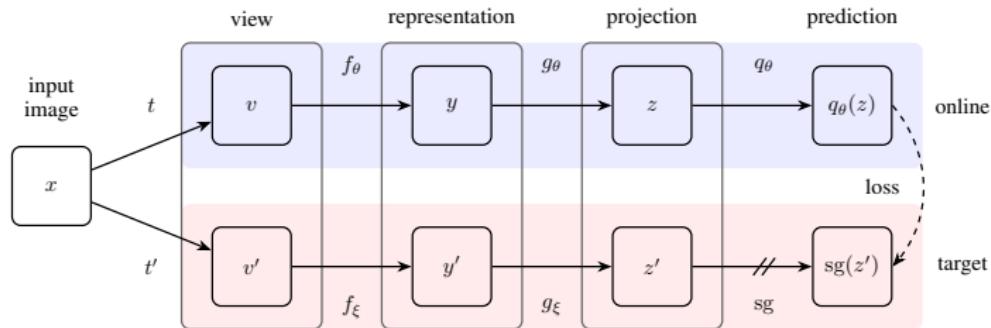
So far: Contrastive loss between exemplar, positive and negatives

- Negative pairs critical (often: large batchsizes, memory banks, custom mining strategies)
- Right augmentation strategy critical

BYOL:

- No negative pair
- No contrastive loss
- More resilient to changes in batch size and set of image augmentations compared to its contrastive counterparts

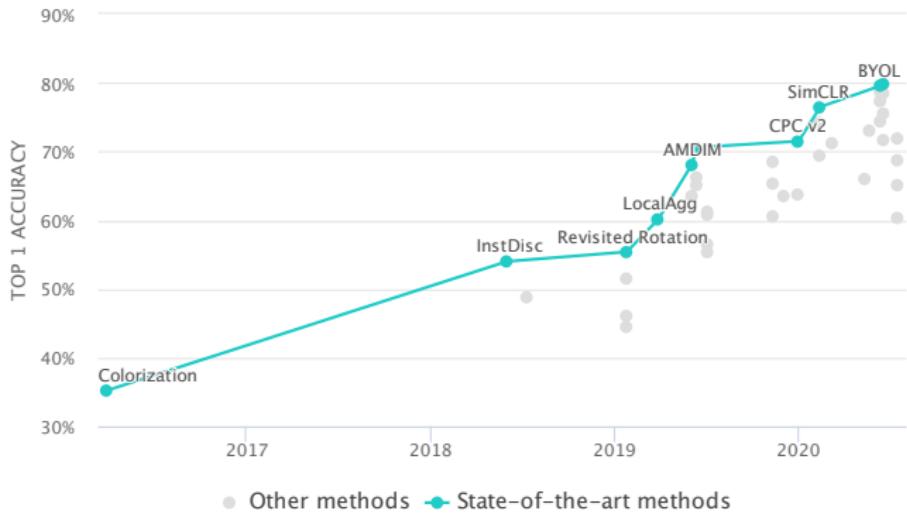
BYOL [34] Method



- Two networks: **online** and **target** network → interact and learn from each other
- In theory: trivial solution possible (e.g. zero for all images)
→ use slow-moving average of the online network as target network
- Loss: MSE of ℓ^2 -normalized predictions (proportional to cosine distance)

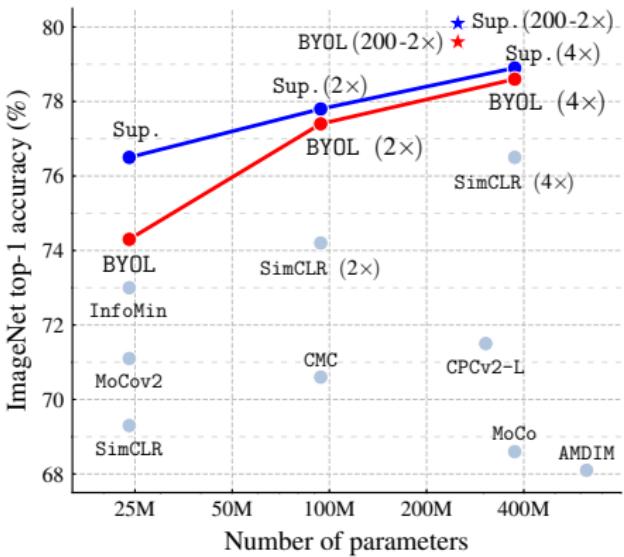
Source: [34]

SSL State of the Art



Source: <https://paperswithcode.com/sota/self-supervised-image-classification-on>

SSL State of the Art (cont.)



Source: [34]

Further Reading

Blogs:

- [https://lilianweng.github.io/lil-log/2019/11/10/
self-supervised-learning.html](https://lilianweng.github.io/lil-log/2019/11/10/self-supervised-learning.html)
- <https://amitness.com/2020/02/illustrated-self-supervised-learning/>
- [https://ankeshanand.com/blog/2020/01/26/
contrative-self-supervised-learning.html](https://ankeshanand.com/blog/2020/01/26/contrative-self-supervised-learning.html)

Others:

- <https://github.com/jason718/awesome-self-supervised-learning>
- <https://www.youtube.com/watch?v=7I0Qt7GALVk>

NEXT TIME
ON DEEP LEARNING

Next Time: Emerging Methods

- Can we process graphs using deep networks?
- Do we really have to learn everything from scratch?
- Let's see whether we can re-use prior knowledge in deep learning...



FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

References



References I

- [1] Özgün Çiçek, Ahmed Abdulkadir, Soeren S Lienkamp, et al. "3d u-net: learning dense volumetric segmentation from sparse annotation". In: [MICCAI](#). Springer. 2016, pp. 424–432.
- [2] Waleed Abdulla. [Mask R-CNN for object detection and instance segmentation on Keras and TensorFlow](#). Accessed: 27.01.2020. 2017.
- [3] Olga Russakovsky, Amy L. Bergman, Vittorio Ferrari, et al. "What's the point: Semantic segmentation with point supervision". In: [CoRR](#) abs/1506.02106 (2015). arXiv: 1506.02106.
- [4] Marius Cordts, Mohamed Omran, Sebastian Ramos, et al. "The Cityscapes Dataset for Semantic Urban Scene Understanding". In: [CoRR](#) abs/1604.01685 (2016). arXiv: 1604.01685.

References II

- [5] Richard O. Duda, Peter E. Hart, and David G. Stork. Pattern classification. 2nd ed. New York: Wiley-Interscience, Nov. 2000.
- [6] Anna Khoreva, Rodrigo Benenson, Jan Hosang, et al. "Simple Does It: Weakly Supervised Instance and Semantic Segmentation". In: arXiv preprint arXiv:1603.07485 (2016).
- [7] Kaiming He, Georgia Gkioxari, Piotr Dollár, et al. "Mask R-CNN". In: CoRR abs/1703.06870 (2017). arXiv: 1703.06870.
- [8] Sangheum Hwang and Hyo-Eun Kim. "Self-Transfer Learning for Weakly Supervised Lesion Localization". In: MICCAI. Springer. 2016, pp. 239–246.
- [9] Maxime Oquab, Léon Bottou, Ivan Laptev, et al. "Is object localization for free? weakly-supervised learning with convolutional neural networks". In: Proc. CVPR. 2015, pp. 685–694.

References III

- [10] Alexander Kolesnikov and Christoph H. Lampert. "Seed, Expand and Constrain: Three Principles for Weakly-Supervised Image Segmentation". In: [CoRR abs/1603.06098](#) (2016). arXiv: [1603.06098](#).
- [11] Tsung-Yi Lin, Michael Maire, Serge J. Belongie, et al. "Microsoft COCO: Common Objects in Context". In: [CoRR abs/1405.0312](#) (2014). arXiv: [1405.0312](#).
- [12] Ramprasaath R. Selvaraju, Abhishek Das, Ramakrishna Vedantam, et al. "Grad-CAM: Why did you say that? Visual Explanations from Deep Networks via Gradient-based Localization". In: [CoRR abs/1610.02391](#) (2016). arXiv: [1610.02391](#).
- [13] K. Simonyan, A. Vedaldi, and A. Zisserman. "Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps". In: [Proc. ICLR \(workshop track\)](#). 2014.

References IV

- [14] Bolei Zhou, Aditya Khosla, Agata Lapedriza, et al. "Learning deep features for discriminative localization". In: Proc. CVPR. 2016, pp. 2921–2929.
- [15] Longlong Jing and Yingli Tian. "Self-supervised Visual Feature Learning with Deep Neural Networks: A Survey". In: arXiv e-prints, arXiv:1902.06162 (Feb. 2019). arXiv: 1902.06162 [cs.CV].
- [16] D. Pathak, P. Krähenbühl, J. Donahue, et al. "Context Encoders: Feature Learning by Inpainting". In:
2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR). 2016, pp. 2536–2544.
- [17] C. Doersch, A. Gupta, and A. A. Efros. "Unsupervised Visual Representation Learning by Context Prediction". In:
2015 IEEE International Conference on Computer Vision (ICCV). Dec. 2015, pp. 1422–1430.

References V

- [18] Mehdi Noroozi and Paolo Favaro. "Unsupervised Learning of Visual Representations by Solving Jigsaw Puzzles". In: Computer Vision – ECCV 2016. Cham: Springer International Publishing, 2016, pp. 69–84.
- [19] Spyros Gidaris, Praveer Singh, and Nikos Komodakis. "Unsupervised Representation Learning by Predicting Image Rotations". In: International Conference on Learning Representations. 2018.
- [20] Mathilde Caron, Piotr Bojanowski, Armand Joulin, et al. "Deep Clustering for Unsupervised Learning of Visual Features". In: Computer Vision – ECCV 2018. Cham: Springer International Publishing, 2018, pp. 139–156.

References VI

- [21] A. Dosovitskiy, P. Fischer, J. T. Springenberg, et al. "Discriminative Unsupervised Feature Learning with Exemplar Convolutional Neural Networks". In: [IEEE Transactions on Pattern Analysis and Machine Intelligence](#) 38.9 (Sept. 2016), pp. 1734–1747.
- [22] V. Christlein, M. Gropp, S. Fiel, et al. "Unsupervised Feature Learning for Writer Identification and Writer Retrieval". In: [2017 14th IAPR International Conference on Document Analysis and Recognition](#) Vol. 01. Nov. 2017, pp. 991–997.
- [23] Z. Ren and Y. J. Lee. "Cross-Domain Self-Supervised Multi-task Feature Learning Using Synthetic Imagery". In: [2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition](#). June 2018, pp. 762–771.

References VII

- [24] Asano YM., Rupprecht C., and Vedaldi A. "Self-labelling via simultaneous clustering and representation learning". In: [International Conference on Learning Representations. 2020.](#)
- [25] Ben Poole, Sherjil Ozair, Aaron Van Den Oord, et al. "On Variational Bounds of Mutual Information". In: [Proceedings of the 36th International Conference on Machine Learning.](#) Vol. 97. Proceedings of Machine Learning Research. Long Beach, California, USA: PMLR, Sept. 2019, pp. 5171–5180.
- [26] R Devon Hjelm, Alex Fedorov, Samuel Lavoie-Marchildon, et al. "Learning deep representations by mutual information estimation and maximization". In: [International Conference on Learning Representations. 2019.](#)

References VIII

- [27] Aaron van den Oord, Yazhe Li, and Oriol Vinyals. "Representation Learning with Contrastive Predictive Coding". In: [arXiv e-prints](#), arXiv:1807.03748 (July 2018). arXiv: 1807.03748 [cs.LG].
- [28] Philip Bachman, R Devon Hjelm, and William Buchwalter. "Learning Representations by Maximizing Mutual Information Across Views". In: [Advances in Neural Information Processing Systems 32](#). Curran Associates, Inc., 2019, pp. 15535–15545.
- [29] Yonglong Tian, Dilip Krishnan, and Phillip Isola. "Contrastive Multiview Coding". In: [arXiv e-prints](#), arXiv:1906.05849 (June 2019), arXiv:1906.05849. arXiv: 1906.05849 [cs.CV].
- [30] Kaiming He, Haoqi Fan, Yuxin Wu, et al. "Momentum Contrast for Unsupervised Visual Representation Learning". In: [arXiv e-prints](#), arXiv:1911.05722 (Nov. 2019). arXiv: 1911.05722 [cs.CV].

References IX

- [31] Ting Chen, Simon Kornblith, Mohammad Norouzi, et al. "A Simple Framework for Contrastive Learning of Visual Representations". In: [arXiv e-prints](#), arXiv:2002.05709 (Feb. 2020), arXiv:2002.05709. arXiv: 2002.05709 [cs.LG].
- [32] Ishan Misra and Laurens van der Maaten. "Self-Supervised Learning of Pretext-Invariant Representations". In: [arXiv e-prints](#), arXiv:1912.01991 (Dec. 2019). arXiv: 1912.01991 [cs.CV].
- [33] Prannay Khosla, Piotr Teterwak, Chen Wang, et al. "Supervised Contrastive Learning". In: [arXiv e-prints](#), arXiv:2004.11362 (Apr. 2020). arXiv: 2004.11362 [cs.LG].

References X

- [34] Jean-Bastien Grill, Florian Strub, Florent Altché, et al. “Bootstrap Your Own Latent: A New Approach to Self-Supervised Learning”. In: [arXiv e-prints](#), arXiv:2006.07733 (June 2020), arXiv:2006.07733. arXiv: 2006 . 07733 [cs.LG].
- [35] Tongzhou Wang and Phillip Isola. “Understanding Contrastive Representation Learning through Alignment and Uniformity on the Hypersphere”. In: [arXiv e-prints](#), arXiv:2005.10242 (May 2020), arXiv:2005.10242. arXiv: 2005 . 10242 [cs.LG].
- [36] Junnan Li, Pan Zhou, Caiming Xiong, et al. “Prototypical Contrastive Learning of Unsupervised Representations”. In: [arXiv e-prints](#), arXiv:2005.04966 (May 2020), arXiv:2005.04966. arXiv: 2005 . 04966 [cs.CV].